

SQL

Structured Query Language

SQL Agenda

— — —

Introduction to SQL

Querying Data

Sorting Data

Filtering Data

Group By

Aggregate Functions

Data Definition Language (DDL)

Constraints

Data Manipulation Language (DML)

Joining Multiple Tables

Conditional Expressions

Filtering Data Part 2

Subquery

Using SET Operators

Data Definition Language (DDL)

Data Types Explained

Depends on DBMS

String, Numeric, Date and Time

Examples

CHAR, VARCHAR, BINARY, BLOB

BIT, BOOLEAN, INT, FLOAT, DOUBLE, DECIMAL

DATE, DATETIME, TIMESTAMP

PostgreSQL Data Types

Boolean (TRUE, FALSE or NULL)

Character - char, varchar, text

Numeric - Integer and Floating-point number

Temporal - Date, Time, Timestamp, Interval

UUID, Array, JSON, hstore, others

CREATE TABLE

CREATE TABLE

Example of Creating the director table

Do it...

CREATE TABLE

(see next slide for details)

Create Tables

— — —

directors

PK director_id
first_name
last_name
Date_of_birth

actors

PK actor_id
first_name
last_name
Date_of_birth

DROP, ALTER,

— — —

```
DROP TABLE IF EXISTS actors;
```

```
ALTER TABLE courses ADD credit_hours INT NOT NULL;
```

TRUNCATE TABLE

Remove all data in a table efficiently and fast

```
TRUNCATE TABLE table_name;
```

```
TRUNCATE TABLE table_name1, table_name2, ...;
```

TRUNCATE TABLE vs DELETE

— — —

Logically the same effect to remove all data with some differences

DELETE - Logs, Allows Rollback

TRUNCATE - No Chance (Exception with a transaction that hasn't committed)

DELETE - Foreign Key OK

TRUNCATE - Not OK (must use delete)

DELETE - Fires Delete Trigger

TRUNCATE - Does Not

DELETE - Can delete partial data

TRUNCATE - Removes all Data

Got it?

— — —

Add TRUNCATE TABLE to directors and actors

Using SQL Constraints

Constraints

— — —

PRIMARY KEY – show you how to define a primary key for a table.

FOREIGN KEY – walk you through the steps of enforcing the relationship between data in two tables using the foreign key constraint.

UNIQUE – ensure the uniqueness of values in a column or a set of columns.

NOT NULL – ensure that the values inserted into or updated to a column are not NULL.

CHECK – validate data before it is stored in one or more columns based on a Boolean expression

Primary Key

Uniquely identify each row in the table

Composite Key - two or more columns

```
CREATE TABLE projects (  
    project_id INT PRIMARY KEY,  
    project_name VARCHAR(255),  
    start_date DATE NOT NULL,  
    end_date DATE NOT NULL  
);
```

	employee_id	course_id	taken_date
	100	3	1987-06-17
	101	3	1989-09-21
	102	3	1993-01-13
	103	3	1990-01-03
	104	3	1991-05-21
	105	3	1997-06-25
	106	3	1998-02-05
	107	3	1999-02-07

Foreign Key

— — —

Link between two tables

```
CREATE TABLE projects (  
    project_id INT AUTO_INCREMENT PRIMARY KEY,  
    ...  
);  
  
CREATE TABLE project_milestones (  
    milestone_id INT AUTO_INCREMENT PRIMARY KEY,  
    project_id INT,  
    FOREIGN KEY (project_id)  
        REFERENCES projects (project_id)  
);
```

Unique

Unique values

```
CREATE TABLE users (  
    user_id INT AUTO_INCREMENT PRIMARY KEY,  
    username VARCHAR(255) NOT NULL UNIQUE,  
    password VARCHAR(255) NOT NULL  
);
```

NOT NULL

Non-NULL values only

```
CREATE TABLE training (  
    employee_id INT,  
    course_id INT,  
    taken_date DATE NOT NULL,  
    PRIMARY KEY (employee_id , course_id)  
);
```

CHECK

Must satisfy a Boolean expression

```
CREATE TABLE products (  
    product_id INT PRIMARY KEY,  
    product_name VARCHAR(255) NOT NULL,  
    selling_price NUMERIC(10,2) CHECK (selling_price > 0)  
);
```

Default Values

Default Values

— — —

- SERIAL
- DEFAULT

```
CREATE TABLE products (  
    product_no SERIAL,  
    price NUMERIC DEFAULT 9.99  
);
```

Challenge



`CREATE TABLE movies`

`(see next slide for details)`

CREATE TABLE movies

movie_id - SERIAL

movie_name - NOT NULL



Challenge



INSERT values into your the movie!

(see next slide for a little test date)

INSERT INTO test data

— — —

Movies

Avatar, 162, 2009-12-18, 1

Star Trek, 127, 2009-05-08, 2

Data Manipulation Language (DML)

INSERT

```
INSERT INTO dependents (  
    first_name, last_name, relationship, employee_id  
) VALUES (  
    'Dustin', 'Johnson', 'Child', 178  
);
```

UPDATE

UPDATE employees

SET last_name = 'Lopez'

WHERE employee_id = 192;

DELETE

— — —

```
DELETE FROM dependents
```

```
WHERE
```

```
    dependent_id = 16;
```

Challenge



`INSERT` values into your tables!

(see next slide for a little test date)

INSERT INTO test data

— — — **Directors**

James Cameron, 1954-08-16

J.J. Abrams, 1966-06-27

Actors

Sam Worthington, 1976-08-02

Zoe Saldana, 1978-06-19

Sigourney Weaver

John Cho

Chris Pine

Database Relationships

What are Database Relationships

One-to-one: A record in one table is related to one record in another table.

One-to-many: A record in one table is related to many records in another table.

Many-to-many: Multiple records in one table are related to multiple records in another table.

Joining Multiple Tables

Joining Multiple Tables


INNER JOINS - Challenge

RIGHT JOINS, LEFT JOINS, FULL JOINS - Challenge


JOINING MORE THAN TWO TABLES - Challenge

UNION, UNION ALL - Challenge


INTERSECT, EXCEPT, Challenge



LEFT INCLUSIVE




LEFT EXCLUSIVE




FULL OUTER INCLUSIVE


SQL JOINS	
LEFT INCLUSIVE SELECT [Select List] FROM TableA A LEFT OUTER JOIN TableB B ON A.Key= B.Key	RIGHT INCLUSIVE SELECT [Select List] FROM TableA A RIGHT OUTER JOIN TableB B ON A.Key= B.Key
LEFT EXCLUSIVE SELECT [Select List] FROM TableA A LEFT OUTER JOIN TableB B ON A.Key= B.Key WHERE B.Key IS NULL	RIGHT EXCLUSIVE SELECT [Select List] FROM TableA A LEFT OUTER JOIN TableB B ON A.Key= B.Key WHERE A.Key IS NULL
FULL OUTER INCLUSIVE SELECT [Select List] FROM TableA A FULL OUTER JOIN TableB B ON A.Key = B.Key	FULL OUTER EXCLUSIVE SELECT [Select List] FROM TableA A FULL OUTER JOIN TableB B ON A.Key = B.Key WHERE A.Key IS NULL OR B.Key IS NULL
INNER JOIN SELECT [Select List] FROM TableA A INNER JOIN TableB B ON A.Key = B.Key	




RIGHT INCLUSIVE



RIGHT EXCLUSIVE



INNER JOIN



FULL OUTER EXCLUSIVE

Joining Multiple Tables

SQL Aliases – make your query shorter and more understandable.

INNER JOIN – introduce you to the join concept and show you how to use the INNER JOIN clause to combine data from multiple tables.

LEFT OUTER JOIN – provide you with another kind of joins that allows you to combine data from multiple tables.

FULL OUTER JOIN – join multiple tables by including rows from both tables whether or not the rows have matching rows from another table.


CROSS JOIN – produce a Cartesian product of rows of the joined tables using the cross join operation.

SELF JOIN – join a table to itself using either the inner join or left join clause.

INNER JOIN

INNER JOIN

— — —



INNER JOIN


SELECT

A.n

FROM A

INNER JOIN B ON

B.n = A.n;



INNER JOIN

— — —

SELECT

A.n

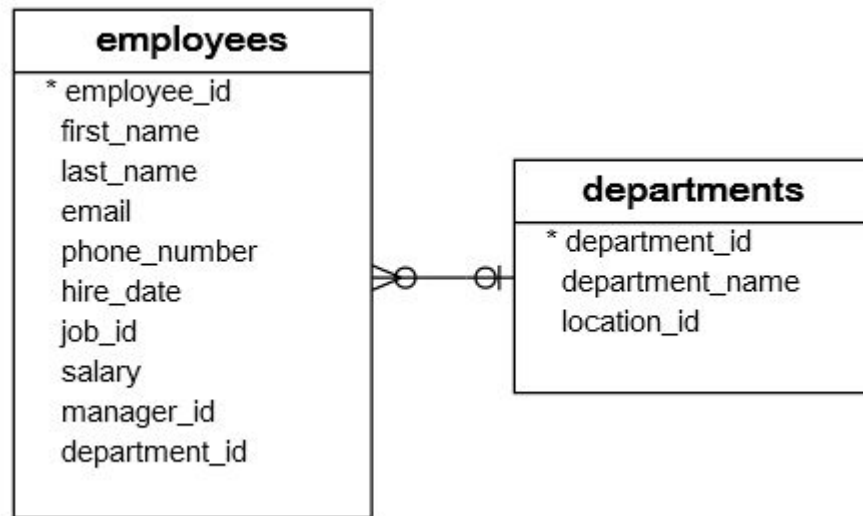
FROM A

INNER JOIN B **ON** B.n = A.n

INNER JOIN C **ON** C.n = A.n;

INNER JOIN

— — —

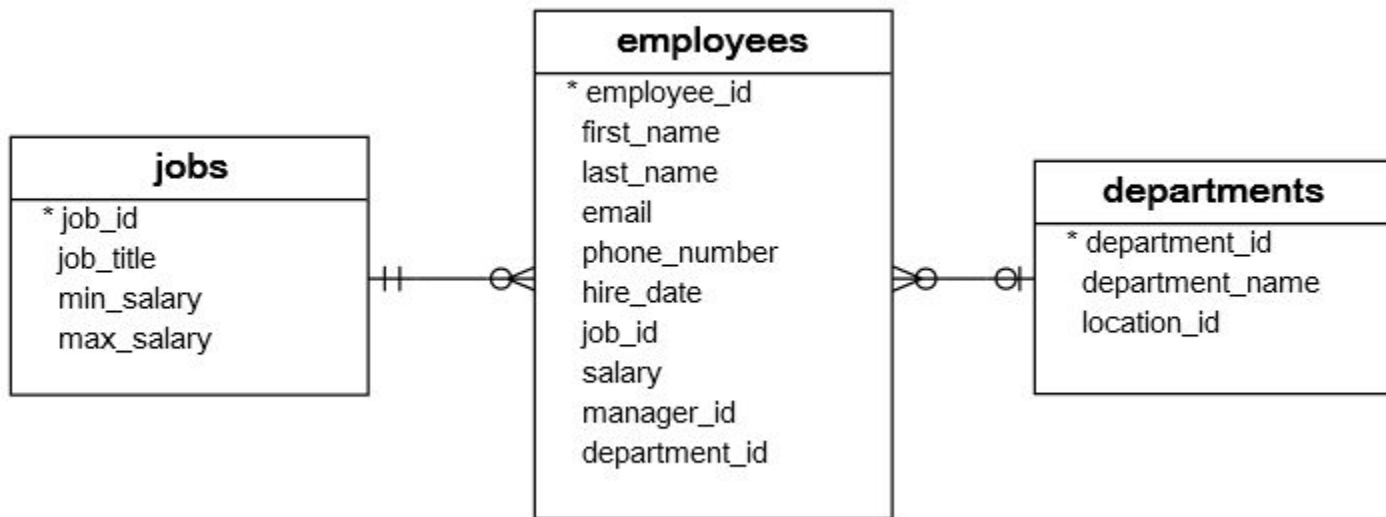


See it...

SQL INNER JOIN

INNER JOIN

— — —



Challenge



Create INNER JOIN for Movie Database

SQL Alias

```
SELECT
    employee_id,
    concat(first_name, ' ', last_name) fullname
FROM
    employees e
INNER JOIN departments d ON d.department_id = e.department_id
```


Diagram annotations:

- A green wavy line underlines the text `fullname` in the `SELECT` clause, with the label "Column alias" written below it.
- A speech bubble points to the text `e` after `employees` in the `FROM` clause, with the label "Table alias" written inside.
- A speech bubble points to the text `d` after `departments` in the `INNER JOIN` clause, with the label "Table alias" written inside.

LEFT, RIGHT, FULL JOIN

LEFT JOIN

LEFT JOIN and LEFT OUTER JOIN



LEFT JOIN


SELECT

A.n

FROM

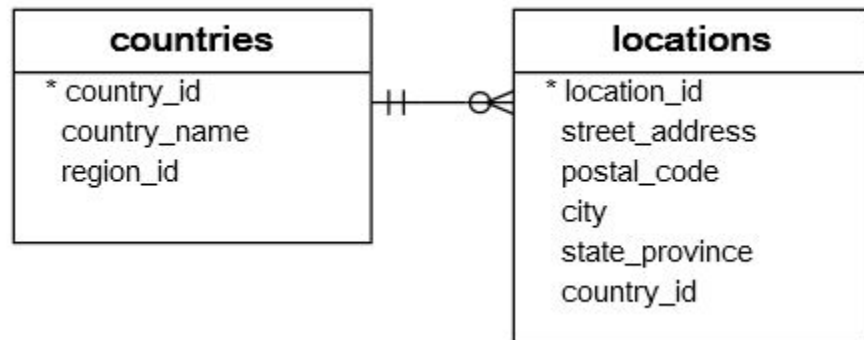
A

LEFT JOIN B ON B.n = A.n;



LEFT JOIN

— — —



See it...

LEFT JOIN

Challenge



Crate LEFT JOIN with Movie Database

Cheat Sheets

Basic Queries

- filter your columns
SELECT col1, col2, col3, ... **FROM** table1
- filter the rows
WHERE col4 = 1 **AND** col5 = 2
- aggregate the data
GROUP by ...
- limit aggregated data
HAVING count(*) > 1
- order of the results
ORDER BY col2

Useful keywords for **SELECTS**:

- DISTINCT** - return unique results
- BETWEEN** a **AND** b - limit the range, the values can be numbers, text, or dates
- LIKE** - pattern search within the column text
- IN** (a, b, c) - check if the value is contained among given.

Data Modification


- update specific data with the **WHERE** clause
UPDATE table1 **SET** col1 = 1 **WHERE** col2 = 2
- insert values manually
INSERT INTO table1 (**ID**, **FIRST_NAME**, **LAST_NAME**)
VALUES (1, 'Rebel', 'Labs');
- or by using the results of a query
INSERT INTO table1 (**ID**, **FIRST_NAME**, **LAST_NAME**)
SELECT id, last_name, first_name **FROM** table2

Views


A **VIEW** is a virtual table, which is a result of a query.
They can be used to create virtual tables of complex queries.

```
CREATE VIEW view1 AS  
SELECT col1, col2  
FROM table1  
WHERE ...
```


The Joy of JOINS



LEFT OUTER JOIN - all rows from table A, even if they do not exist in table B



INNER JOIN - fetch the results that exist in both tables



RIGHT OUTER JOIN - all rows from table B, even if they do not exist in table A

Updates on JOINed Queries

You can use **JOINS** in your **UPDATES**:

```
UPDATE t1 SET a = 1  
FROM table1 t1 JOIN table2 t2 ON t1.id = t2.t1_id  
WHERE t1.col1 = 0 AND t2.col2 IS NULL;
```

NB! Use database specific syntax, it might be faster!

Semi JOINS

You can use subqueries instead of **JOINS**:

```
SELECT col1, col2 FROM table1 WHERE id IN  
(SELECT t1_id FROM table2 WHERE date >  
CURRENT_TIMESTAMP)
```

Indexes

If you query by a column, index it!
CREATE INDEX index1 **ON** table1 (col1)

Don't forget:

- Avoid overlapping indexes
- Avoid indexing on too many columns
- Indexes can speed up **DELETE** and **UPDATE** operations

Useful Utility Functions

- convert strings to dates:
TO_DATE (Oracle, PostgreSQL), **STR_TO_DATE** (MySQL)
- return the first non-NULL argument:
COALESCE (col1, col2, "default value")
- return current time:
CURRENT_TIMESTAMP
- compute set operations on two result sets
SELECT col1, col2 **FROM** table1
UNION / EXCEPT / INTERSECT
SELECT col3, col4 **FROM** table2;

- Union* - returns data from both queries
- Except* - rows from the first query that are not present in the second query
- Intersect* - rows that are returned from both queries

Reporting

Use aggregation functions

- COUNT** - return the number of rows
- SUM** - cumulate the values
- AVG** - return the average for the group
- MIN / MAX** - smallest / largest value